

RED-PD: RED with Preferential Dropping

<http://www.aciri.org/red-pd>

Ratul Mahajan


Sally Floyd

David Wetherall

Background

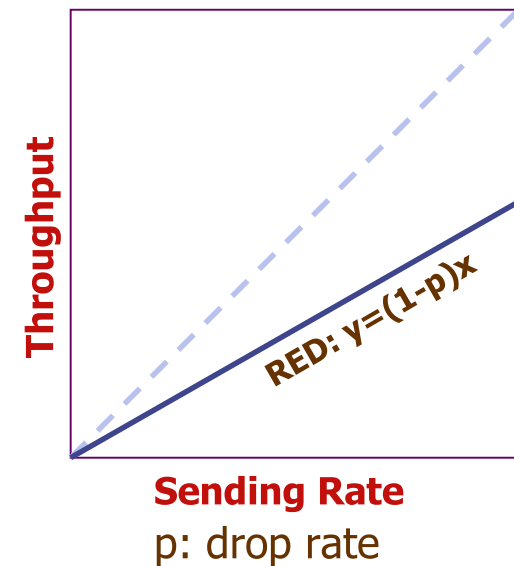
- ❖ FIFO – first in first out
 - drop tail

- ❖ RED – Random Early Detection
 - detect incipient congestion
 - distribute drops fairly



Problem

- ❖ High bandwidth flows increase the drop rate at the router
- ❖ Without flow based differentiation all flows see the same drop rate
 - flows get more by sending more



Solution: router mechanisms to protect rest of the traffic from high bandwidth flows

Relevance

- ❖ Where is congestion?
 - Backbone (?)
 - Edge routers
 - Exchange points
 - Intercontinental links

- ❖ Starting point for aggregate based congestion

A Possible Approach – Per-flow state

Examples: FQ, DRR, CSFQ, FRED

Sequester all flows from each other

- + provide full max-min fairness
 - o required for best-effort traffic?
- state for ALL the flows passing through the router, most of which are small “Web mice”

RED-PD's Approach – Partial Flow State


State for high bandwidth flows only

- ❖ Identify high bandwidth flows during times of congestion. These are called monitored flows.
- ❖ Preferentially drop from monitored flows to limit their throughput.

Combines the simplicity of FIFO with the protection of full max-min fair techniques that keep per-flow state.

Why Partial Flow State Approach Works

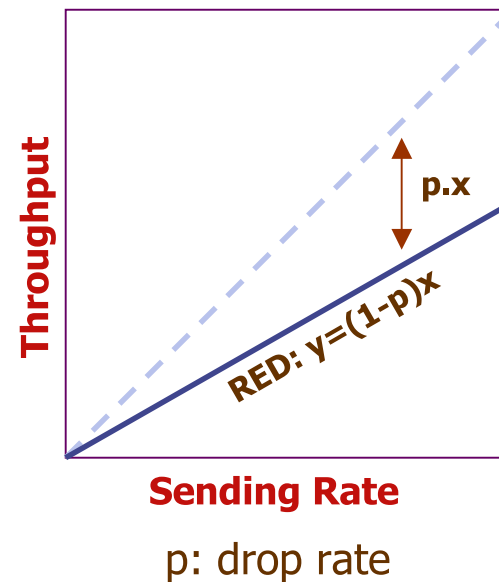
What fraction of flows get what fraction of bytes over different time windows.



Bandwidth distribution is very skewed - a small fraction of flows accounts for most of the bandwidth.

Identification Approach – Drop History

Flows that send more
suffer more drops




- ❖ Cheap means for identification
- ❖ Drop history contains flows that have been sent congestion signal

Defining “High Bandwidth”

- ❖ Pick a round trip time (RTT) R
 - call a TCP with RTT R as reference TCP
- ❖ High bandwidth flows:
 - All flows sending more than the *reference* TCP flow
- ❖ Target bandwidth $Target(R,p)$ given by the TCP response function
 - $Target(R,p) \approx \frac{\sqrt{1.5}}{R \sqrt{p}}$ packets/second
where p is the drop rate at the output queue

Identifying High Bandwidth Flows



- ❖ TCP suffers one drop in a congestion epoch
- ❖ Length of congestion epoch can be computed using drop rate
- ❖ $CELength(R,p) = \frac{R}{\sqrt{1.5} \sqrt{p}}$ congestion epoch length of the reference TCP
 - identify flows with one or more drops in $CELength(R,p)$ seconds

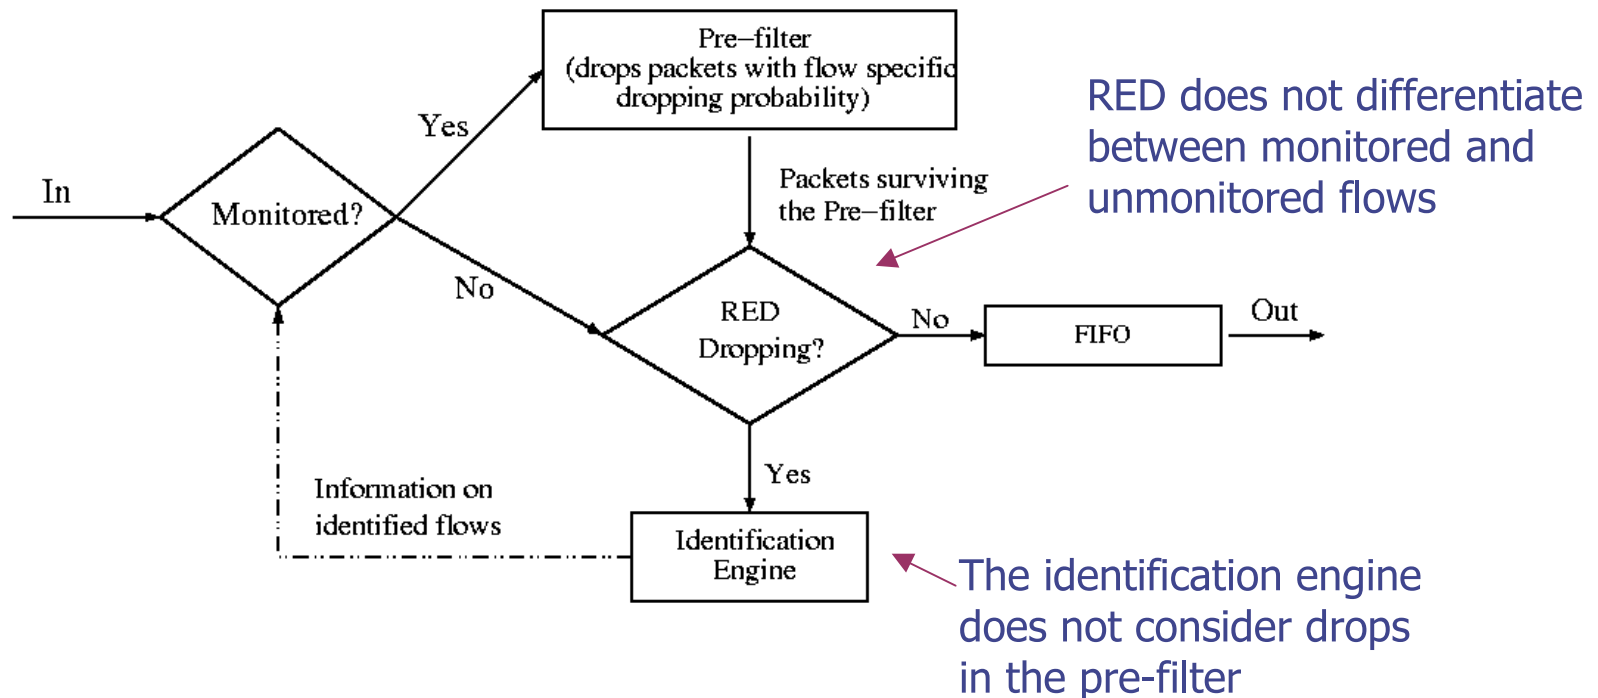
Tuning Identification

1. Issue: all flows suffer a drop once in a while
 - keep the drop history of $K.CELength(R,p)$ seconds and identify flows with K or more drops ($K>1$)
2. Issue: multiple losses in a window of data
 - maintain the drop history as M separate lists instead of a single list
 - length of each list is $(K/M).CELength(R,p)$ seconds
 - identify flows with drops in K or more of the M lists
 - lets go flows with more than K drops in a short period

Preferential Dropping

- ❖ Probabilistically drop packets from monitored flows before they enter the output queue
- ❖ The dropping probability is computed to bring down rate of a monitored flow into the output queue to $Target(R,p)$
 - $probability = 1 - Target(R,p)/arrival\ rate$

Architecture



- ❖ If a monitored flow is identified again, the pre-filter dropping probability is too small
- ❖ If a monitored flow is suffering too few drops, the pre-filter dropping probability is too large


Probability Computation

- ❖ Decrease dropping probability of a monitored flow with no drops in M lists by halving it
 - upper bound on reduction in one round
- ❖ Increase dropping probability of an identified flow
 - an identified flow could be either a monitored flow or a newly identified flow.
- ❖ Don't change dropping probability too fast

Probability Increase

- ❖ Add an increase quantum to existing probability
- ❖ The increase quantum should depend on
 - *ambient* drop rate (drop rate at the output queue)
 - sending rate of a flow
- ❖ $\text{increase quantum} = (\text{number_drops}/\text{avg_drops}).p$
 - upper bound on increase in one step

Effect of RED-PD



p : drop rate with RED
 p' : ambient drop rate
with RED-PD


Target Bandwidth

- ❖ Reduction in ambient drop rate ($p' < p$)
- ❖ Full max-min fair in the extreme case ($p' = 0$)

Evaluation

- ❖ Fairness
- ❖ Response time
- ❖ Effect of target RTT R
- ❖ Probability of identification
- ❖ Persistent congestion throughput
- ❖ Web traffic
- ❖ Multiple congested links
- ❖ TFRC
- ❖ Byte mode operation

Fairness



10 Mbps link

R= 40ms

Flow Index

1-3 30ms TCP

4-6 50ms TCP

7-9 70ms TCP


10 5 Mbps CBR

11 3 Mbps CBR


12 1 Mbps CBR

Iterative probability
changes successfully
approximate fairness

Response Time



10 Mbps link
1 CBR flow
9 TCP flows
 $R = 40$ ms



0.25Mbps

4Mbps

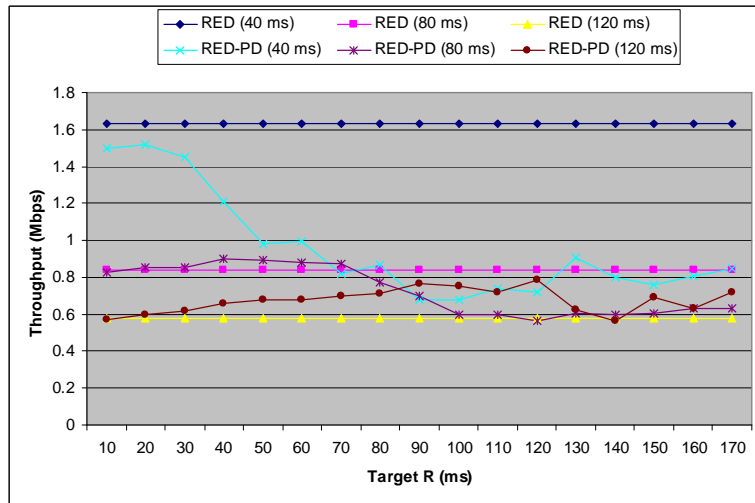
0.25Mbps

Sending rate of CBR changes with time

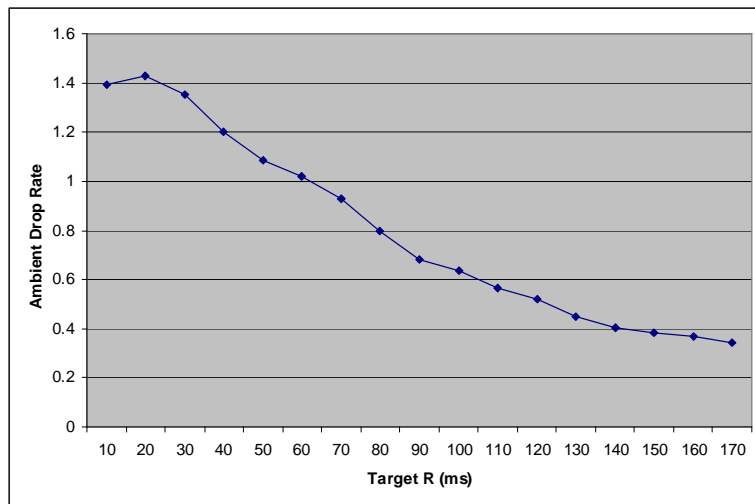
Flow brought down in 6s
and released in 5s

The speed of RED-PD's reaction depends on ambient drop rate and flow's sending rate. Its faster when either is higher

Effect of target RTT R



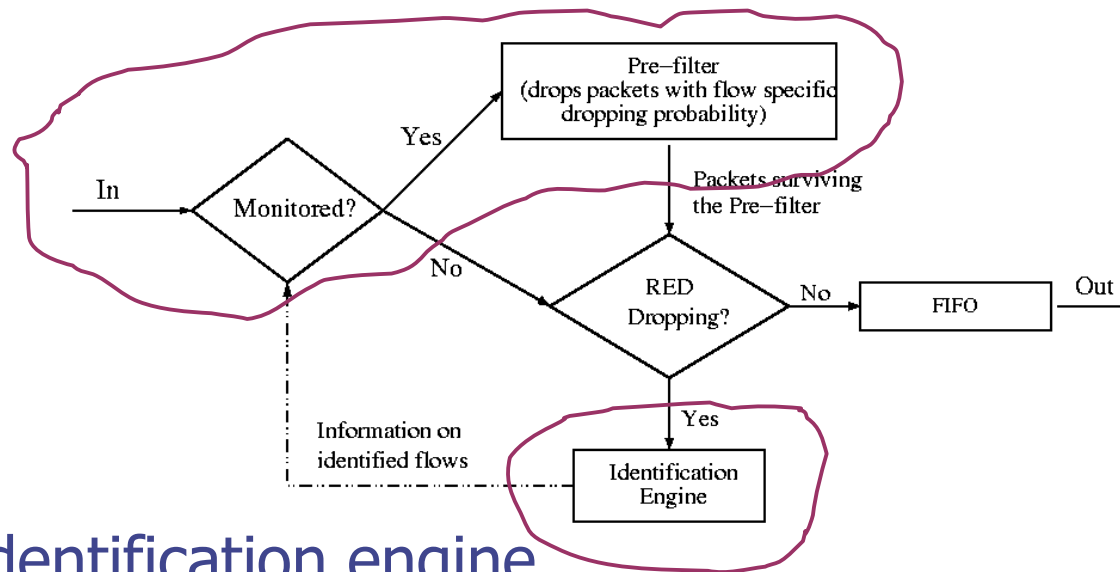
10 Mbps link
14 TCP flows
2 each of RTT 40, 80 & 120 ms
and 8 of RTT 160 ms



Increasing R

- increases fairness
- increases state
- decreases ambient drop rate

Implementation Feasibility



- ❖ Identification engine
 - state for drop history
 - not in fast forwarding path
- ❖ Flow classification and pre-filter
 - hash table of all flows being monitored
 - drop with the computed probability

Comparison

	State for what?	What state?	Fast-path processing?	When processing?
FQ/DRR	All flows	Queues	Queue management and scheduling	Arrival/Departure
FRED	All buffered flows	Number of buffered packets	Drop probability computation and coin tossing	Arrival/Departure
CSFQ	All flows	Arrival rate estimate, time of last packet	Update arrival rate estimate, change header	Arrival
RED-PD	High bandwidth flows	Dropping probability, drop history	Coin tossing	Arrival

Future Work

- ❖ State requirements
- ❖ Dynamically varying R
 - Target ambient drop rate
 - State limitations
- ❖ Misbehaving flows

Conclusions

- ❖ Increased fairness using much less state
- ❖ Tunable degree of fairness
- ❖ Reasonable response time
- ❖ Lightweight and easy to implement

A continuum of policies

