

Eat all you can in an all-you-can-eat buffet: A case for aggressive resource usage

Ratul Mahajan


Jitu Padhye, Ramya Raghavendra, Brian Zill

Microsoft Research

Avoiding hunger

Alice walks into a restaurant

**All you can eat
BUFFET
\$6.95 !**



How much to eat to minimize the chance of getting hungry before the next meal?

1. As much as stomach space allows; or
2. Based on expected time until next meal

Avoiding packet loss

Bob needs to transmit data over a noisy wireless channel


How much FEC to add to minimize the chance of losing packets?

1. As much as the available spectrum; or
2. Based on expected bit error rate

The focus on efficiency in current designs

Operating at the sweet spot tends to be the goal

- Low efficiency beyond it




Inappropriate in many scenarios

- If the resource is of “use it or lose it” kind
- If the sweet spot is hard to determine

The Buffet principle

*Continue using more resources as long as the marginal cost **can be driven** lower than the marginal benefit*



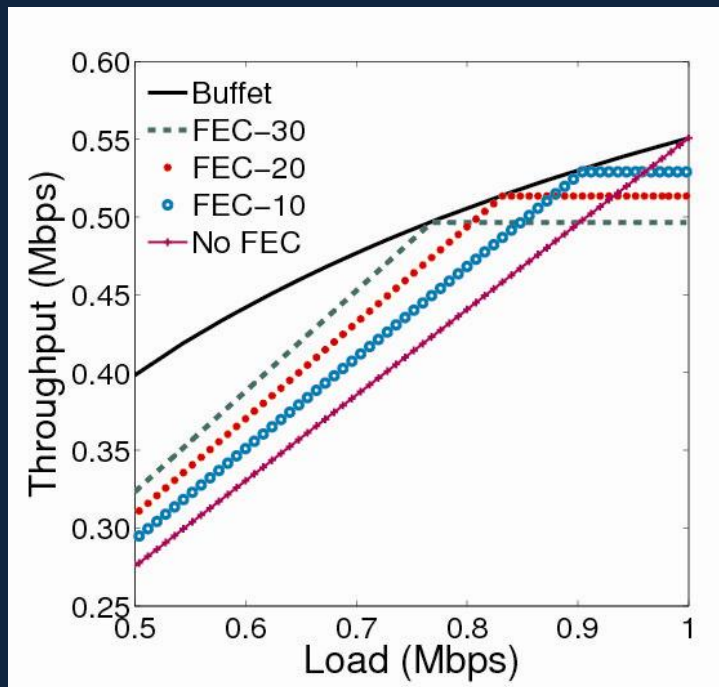
Key strength: performance limited by available resources, not artificial design choices

Key challenge: the default way to aggressively use resources is often problematic

Case: Adding FEC to data transmissions

Current practice: # of FEC bits is independent of load and available spectrum

Buffet: aggressively use available spectrum for FEC bits



Simulation results with 1000-byte packets, 1 Mbps channel, 10^{-6} BER


Making Buffet FEC practical

Greedy addition of FEC can hurt
with multiple transmitters

- Especially for CSMA systems (e.g., WiFi)

Possible solution [under development]

- Embed bits in separate packets
- Send at low priority
- Keep FEC packets small



Data


FEC

Low priority FEC

Case: Erasure coding for lossy paths

Current practice: balance # of coded packets and bandwidth cost

Buffer: aggressively uses all available capacity



Managing the impact of aggressive redundancy

- Use priority on packets if router support is available
- Transmit only when queue is empty [under development]

More case studies

Mobility updates in ad hoc networks

- Must balance freshness of info and bw overhead
- Buffet: Aggressively send updates per spare capacity
 - Prioritization can minimize impact on data transfers

Routing in delay-tolerant networks (DTNs)

- Must balance message delivery prob. and fairness
- Recent design (RAPID) uses the Buffet principle
 - A utility-driven framework prevents disproportionate resource usage by some messages

Yet more case studies

Using replication to boost reliability

- Across disks within a computer
 - Use background tasks to manage overhead
- Across multiple computers
 - Use background transfers to manage overhead

Pre-loading binaries into memory

Pre-fetching Web pages into caches

Speculative execution of program branches

Considerations in applying the Buffet principle

What are the challenges?

What resources does it apply to (naturally)?

What are the side-effects?

Challenges in applying the principle

Aggressive resource usage should not detract from productive work

- **Helpful techniques:** explicit or implicit priorities; opportunistic usage; and utility-driven usage

Quantifying the marginal cost and benefit

- Must include the impact of side-effects as well
- Precise accounting not needed in many situations

Applicable resources

More natural for non-conservable resources

- E.g., storage, bandwidth, computation
- Feasible for others as well (e.g., battery)

Easier when the resource is not shared with
non-Buffer users

Side-effects of aggressive resource use

New bottlenecks or concerns

- E.g., disk I/O bw, energy consumption

Increased hardware wear-and-tear

Task completion latency may increase

Performance is now coupled with load

Conclusions

*Continue using more resources as long as the marginal cost **can be driven** lower than the marginal benefit*

The Buffet principle advocates aggressive resource usage instead of a singular focus on efficiency

It has the potential to provide the best performance for the level of available resources